



HITHERTO

2026

Global Dialogue on AI Governance

Sovereignty as an Architectural Choice

Prepared by: R. Verity Vabolis - Director - Hitherto AI, Flux Hall
Academy & Fathom Lab

Sovereignty as an Architectural Choice

Replacing policy-heavy 'finger-wagging' with the technical teeth of architectural enforcement to secure global sovereignty, clinical integrity, and shared reciprocity in the AI era.

Geneva, July 2026

Position Paper for the Global Dialogue on Artificial Intelligence Governance

I. Executive Summary

Current international AI governance efforts frequently stall at the level of high-level ethics and "policy promises." This position paper advocates for a paradigm shift toward Architectural Enforcement, the embedding of safety, privacy, and sovereignty directly into the technical fabric of AI infrastructure. By moving away from subjective "finger-wagging" and toward objective, measurable technical controls, we can ensure that the rapid evolution of digital technologies supports human rights and international law by default.

II. What is a Successful Dialogue?

To make the first Global Dialogue a success, the outcome must move beyond "governance by vibes" and toward Architectural Enforcement. If we leave Geneva with only a list of high-level principles and no technical blueprints, we haven't governed; we've merely "cosplayed."

A successful dialogue would achieve three concrete outcomes:

Mandatory Technical Standards for High-Risk Domains:

The Dialogue must transition from "Policy Promises" to identifying specific "priority actions". Success means establishing an international commitment to Partitioned Architecture in life-critical sectors like medicine. We must decouple "Medical Truth" from the hallucinations of General Knowledge Models (GKMs) currently used by over 400,000 physicians.

Validation of Stateless Infrastructure: A key "practical lesson" would be the adoption of "Stateless-by-Design" protocols, such as the Bauta Valve.

By making the retention of patient/citizen PII mathematically impossible, we move privacy from a legal debate to a technical certainty.

Sustainable Modernization for the Global South: To truly "bridge AI divides," the Dialogue must validate data-sovereignty models. Success looks like a global framework for anonymized data-licensing, allowing developing nations to fund their own infrastructure rather than relying on external aid.

Finally, the Dialogue must mandate Technical Literacy Minimums (TLM) for oversight bodies. "Robust human oversight" is an illusion if the humans in charge cannot define or verify the "measurable thresholds" We advocate for the mandatory integration of Deterministic Remediation Protocols and Multi-Temporal Performance Baselines as non-negotiable architectural prerequisites. These categories of Technical Recourse ensure that mission-critical operations are governed by measurable integrity gates rather than subjective ethics checklists.

III. Narrowing Priorities

In Safe AI and Human Oversight, we advocate for Technical Literacy Minimums (TLM) for the humans tasked with governance. Oversight is an illusion without competence; therefore, we emphasize that governance professionals must be equipped to identify and verify measurable controls like Automated System Integrity Gates, Deterministic Remediation Protocols, and Multi-Temporal Performance Baselines instead of finger pointing and report writing.

Through Flux Hall Academy, we address Capacity-building by modernizing the delivery of technical education. We advocate for a global shift toward high-engagement, domain-specific learning frameworks that respect the expertise of adult learners. This model moves beyond basic vocabulary to actual system mastery, enabling local stakeholders to manage their own Partitioned/Moded Intelligence systems.

Regarding Societal Implications, we focus on bridging the "Integrity Gap" in high-risk domains like medicine. By deploying Stateless-by-Design infrastructure like the Bauta Valve, we ensure that "Medical Truth" is decoupled from the hallucinations of general-purpose engines, protecting both patient safety and cultural context.

Finally, we promote Digital Sovereignty for the Global South. By combining modernized adult-learning frameworks with anonymized data-licensing, nations can fund and operate their own infrastructure. Our goal is to provide a scalable blueprint for Inclusive Modernization where international cooperation is based on shared technical competence and standard-driven enforcement.

IV. Emerging Issues at Hand

We identify five critical cross-cutting issues that require urgent integration into the Global Dialogue:

1. Minimum Standards for Public-Facing Technology: Current governance is reactive. We must establish "pre-deployment minimums" that ensure any technology reaching the public has passed rigorous, non-optional safety thresholds. Safety cannot be a "feature" added post-launch; it must be an architectural prerequisite.

2. Standards for Certifying Bodies: To eliminate "governance by vibes," we must establish meaningful minimum standards for the entities that certify technology and professionals. A certification is meaningless if the "expert" cannot demonstrate a post-audit methodology or identify technical controls.

We propose a move toward Competency-Based Certification for governance practitioners.

3. Diverse Technical Evolution: We must move beyond "Data Gluttony." The current paradigm, focused almost exclusively on Transformers and neural networks, is not the "end-all" of intelligence. Governance must protect and fund research into alternative architectures that do not rely on mass extraction of earth's resources or data. Humanity cannot make progress if progressive, efficient research is suppressed by the dominance of compute-heavy models.

4. Intentional Scaling (Safety Over Speed): We advocate for the "Everything-By-Default" (EBD) standard, where safety, privacy, and auditability are non-negotiable defaults rather than optional toggles. This involves a conscious choice to limit speed and scale in favor of clinical and societal safety.

5. Finally, we propose the introduction of Neutral Party Liaisons, technical experts who facilitate high-friction cooperation between ministries or nations. These parties focus solely on the science and handling of tech, ensuring that geopolitical friction doesn't stall modernization or safety implementation.

V. Impact of Governance Proper

Significant Challenge: The Clinical Integrity Gap. > Currently, over 400,000 physicians utilize consumer-grade General Knowledge Models (GKMs) for life-critical diagnostics. Research (Stanford/Microsoft) reveals accuracy as low as 34-41% in complex clinical reasoning, despite high USMLE scores. This "Integrity Gap" is a structural failure of current governance: we are treating medical decision-support as a consumer convenience rather than Digital Public Infrastructure (DPI). Furthermore, the governance field suffers from "competence-free" certification, professionals who understand the vocabulary of "risk appetite" but lack the technical literacy to implement a post-audit rollback trigger or measurable control.

Significant Opportunity: Architectural Enforcement. > The most significant advance is the shift toward Architectural Enforcement through Stateless-by-Design protocols and Partitioned Intelligence. By using infrastructure like the Bauta Valve, we can move safety from a policy promise to a technical mandate, ensuring zero-persistence of PII and decoupling "Medical Truth" from generalist hallucinations.

Inclusive Modernization and Sovereignty: > In developing regions, specifically Africa, there is a transformative opportunity for Digital Sovereignty. Rather than relying on external aid or debt, nations can fund their own infrastructure by anonymizing and licensing data. By applying andragogical (adult-learning) frameworks through Flux Hall Academy, we empower local oversight bodies to meet Technical Literacy Minimums. This ensures that "international cooperation" is facilitated by neutral parties focused on science and handling, rather than geopolitical friction. We advocate for an "Everything-By-Default" (EBD) standard where safety and auditability are non-negotiable architectural defaults. We're not proposing that we address the world's problems, but we are proposing that world address the problems by similar standards. In 2026 anything less is an indication of carelessness, laziness, or greed.

VI. Impact of Governance Proper

The Global Dialogue must transition international cooperation from high-level rhetoric to Architectural Enforcement. Its most vital role is to catalyze a global shift toward Sovereignty and Autonomy, ensuring that those currently perceived as "behind" can reach parity with those perceived as "ahead".

To achieve this, the Dialogue must facilitate three practical pillars:

Institutionalizing Neutral Party Liaisons: By deploying technical intermediaries to reduce friction between nations, the Dialogue ensures that modernization continues based on science and handling, rather than geopolitical agendas. This provides the "connective tissue" for a multipolar technological landscape.

Standardizing Technical Literacy Minimums (TLM): Robust human oversight is an illusion without competence. The Dialogue should establish minimum competency standards that empower oversight bodies to verify measurable controls.

Enabling Digital Sovereignty: By promoting models such as anonymized data-licensing and Partitioned Intelligence, the Dialogue allows nations to fund and operate their own infrastructure.

Ultimately, a successful Dialogue rebalances global power. By establishing an "Everything-By-Default" (EBD) standard for safety and auditability, we reduce dependency on a few dominant interests. If the world adopts better, safer, and more autonomous standards, nations rise together in a collective ecosystem. In this future, the interests currently "In Control" are no longer hegemons but participants in a superior, shared global framework. The result is a world where Sovereignty is the default, and technological progress is no longer synonymous with extraction or dependency.

VII. The Added Value of the Dialogue

The Dialogue should build upon the foundational work of the **ITU AI for Good Global Summit** and the **Independent International Scientific Panel on AI**. It must align with the mandate of **General Assembly resolution A/RES/79/325** to connect existing efforts and identify practical lessons.

The unique added value of this Dialogue is its capacity to institutionalize **technological systems of reciprocity**. By utilizing **Neutral Party Liaisons** to facilitate these systems, the Dialogue can move the global community beyond passive observation toward active, secure data-intelligence exchange. In this framework, one country's data trains another's domain-specific intelligence; consequently, each nation gains a tangible stake in the safety, security, and overall success of the other's systems.

Currently, critical tools like medical decision support are treated as consumer conveniences. The Dialogue's added value lies in redefining these as **Digital Public Infrastructure (DPI)**. By adopting **Stateless-by-Design protocols**, such as the Bauta Valve, the Dialogue ensures that this reciprocity is mathematically secure, making patient/citizen PII leakage impossible while promoting interoperability and the trust of the respective peoples of nations engages in these initiatives and partnerships.

Furthermore, by connecting these technical standards with modernized, andragogical technical education (the Flux Hall model), the Dialogue ensures that "Human Oversight" is backed by the competence required to manage these systems. This shifts the global dynamic from hegemonic dependency to a network of autonomous, technically-literate peers. This ensures that as countries rise, they do so together, underpinned by shared, high-fidelity standards rather than mere "finger-wagging" with no recourse.

VIII. Multi-Stakeholder Considerations

We recommend moving away from "Consensus-by-Exhaustion" models toward **Multi-Objective Decision Analysis (MODA)** and **Multi-Actor Multi-Criteria Analysis (MAMCA)**. The current four-tiered "funnel" structure often filters out minority technical truths in favor of high-level jargon. Instead, the Dialogue should employ a **Grid Model of Technical Reciprocity**.

1. Mapping Friction via MODA: Rather than seeking a single "agreed-upon" recommendation, the Dialogue should produce a Sensitivity Map. This map visualizes the mathematical trade-offs between speed, scale, and safety. By weight-tracking agreements without betraying the individual participants, MODA allows "ground truth," such as the need for Technical Recourse in post-audit protocols, to remain visible even when it conflicts with the aggressive scaling interests of dominant actors.

2. Neutral Party Synthesis: We propose the use of Neutral Party Liaisons, technical intermediaries focused strictly on science and handling, to moderate breakouts. These parties facilitate systems where one nation's data can train another's intelligence securely. This creates a vested interest in mutual success, ensuring that the "logical and reasonable" majority is not silenced by a "bullish" few.

3. Everything-By-Default (EBD) Sandboxes: Structure thematic sessions as "stress-test sandboxes" where proposed governance approaches are measured against a baseline of Architectural Enforcement. If a proposal cannot meet the EBD floor for auditability and clinical safety, its "weight" in the final summary is mathematically adjusted. This ensures that the Dialogue results in a Roadmap of Recourse, providing a technical and logical baseline that empowers the sovereign majority and prevents the "cosplaying" of governance.

IX. Underrepresented Voices

We identify three critical groups currently muffled by the "policy funnel":

Ground-Truth Practitioners: For example, the 400,000+ physicians and subspecialists currently utilizing consumer-grade AI for diagnostics. Their clinical reality, where accuracy drops to 34–41% in complex reasoning, is rarely represented in "high-level" ethics discussions.

Localized Sovereignty Developers: Independent innovators from the Global South (e.g., participants of the 2026 India AI Impact Summit) who are building domain-specific models for agriculture and medicine.

Their perspective on Digital Sovereignty is vital to counteracting the current "Data Gluttony" paradigm.

The Distrustful Citizenry: In regions where up to 80% of the population does not trust their government with personal data, the perspective of the "unrepresented end-user" is paramount.

To include them, we must move beyond symbolic representation. We propose a Ground Truth Audit where these stakeholders are invited to stress-test proposed governance frameworks against their specific technical and ethical frictions. Their input should be weighted using MODA (Multi-Objective Decision Analysis) to ensure that the "bullish" interests of dominant actors do not drown out the essential safety requirements of the clinical and localized majority. Inclusion must be an architectural requirement, not a side event.

X. Innovative Engagement

We propose two innovative formats to foster genuine technical reciprocity:

1. Digital Delegate Simulations: Each participating entity, Government or Stakeholder, can utilize Open Source models to train a "Digital Delegate" on their specific policy constraints, cultural nuances, and technical requirements. These delegates can engage in high-speed, background negotiations during the Dialogue, surfacing areas of alignment that might take months for human counterparts to identify. This format respects the time of adult experts while allowing for an exhaustive exploration of technical interoperability.

2. Real-Time Alignment Visualization (The Reciprocity Grid): Using MODA tools, the Dialogue should provide real-time visualization of shared weights and agreements. This allows a small, developing nation to see that their technical safety standards are 80–100% aligned with those of a major power, even if they remain anonymous. This "Reciprocity Grid" builds Citizen Trust by proving that logic and "Ground Truth" create bridges where geopolitical names might create walls.

3. Protocol War-Gaming: Rather than panels, we suggest "Sandboxes" where proposed governance standards, like the Bauta Valve, are live-tested. Participants are tasked with finding PII leaks or clinical hallucinations within a stateless architecture. This shifts the focus from "Policy Promises" to Architectural Enforcement, ensuring that only those standards that can survive a technical "Ground Truth Audit" are included in the Co-Chairs' summary.

XI. Concrete Solutions to Address the Challenges

We advocate for moving beyond "Soft Governance" toward Architectural Enforcement through six scalable practices and platforms:

Stateless-by-Design Infrastructure: This privacy and security protocol mandates a total memory-wipe between data batches. By ensuring Zero-Persistence, the retention or leakage of patient/citizen PII (Personally Identifiable Information) becomes mathematically impossible, moving privacy from a policy promise to a technical mandate.

Partitioned/Moded Intelligence (Arête Techne): To bridge the "Domain Integrity Gap" where generalist engines show accuracy as low as 34% in Life Critical domains, we utilize strictly partitioned, mission-specific architecture. Our Fieldwise training protocols on moded/partitioned systems ensure fidelity in practice.

Technical Literacy Minimums (TLM): "Robust human oversight" is an illusion without competence. We employ an andragogical framework to certify oversight bodies in identifying measurable technical controls, such as , Deterministic Remediation Protocols, and Automated System Integrity Gates, are results of a technical audit rather than "finger-wagging."

Legacy Modernization: The framework represents a multi-phase systems engineering strategy designed to transition fragmented, heterogeneous infrastructure into a unified "Living Lab". Rather than relying on isolated point-solutions, the framework establishes a "Secure by Default" foundation, utilizing Zero Trust Network Architecture and a unified data ontology, to ensure every technological addition is interoperable and sustainable from the outset.

This comprehensive modernization is governed by a Portfolio Integration Protocol, which mandates that all new technologies pass a rigorous rubric for security, interoperability, and measurable evidence of risk reduction. By integrating ambient intelligence inhabitants for proactive safety nets and secure mentorship layers for real-time workforce upskilling, the framework systematically designs out human error and systemic failure.

Ultimately, this process converts a legacy environment into a "High-Fidelity Validation Dataset," providing the ultimate proof-of-work for technology-enabled, safe and sovereign mission-critical operations.

Sovereign Reciprocity Mechanisms: We propose the use of Neutral Party Liaisons to facilitate anonymized data-licensing. This empowers nations to fund their own infrastructure while establishing a "Scientific Commons" where one country's data securely trains another's intelligence.

By establishing an **Everything-By-Default (EBD) standard**, we ensure that safety, privacy, and clinical/domain accuracy are non-negotiable architectural prerequisites for any public-facing technology.